

# Continuously Tracking Core Items in Data Streams with Probabilistic Decays


Junzhou Zhao Pinghui Wang Jing Tao Shuo Zhang John C.S. Lui

Xi'an Jiaotong University  
The Chinese University of Hong Kong




# Expensive to Process Big Data Streams


- Data streams are ubiquitous:
  - email stream, tweets stream, news stream, etc
  - geo-location stream generated by taxis, IoT devices, LBSNs, etc
  - user consuming record stream from Amazon, Taobao, etc
- Applications:
  - real-time trending topic detection
  - network security monitoring
  - online collaborative filtering
- However, the **high speed** and **large volume** cause troubles.



# Two Ways for Handling Big Data Streams




- scale up computation power




- reduce data complexity
  - cheap and green
  - rely on clever algorithms

# Two Ways for Handling Big Data Streams



- scale up computation power



this work



- reduce data complexity
  - cheap and green
  - rely on clever algorithms


# Need for Reducing Data Complexity

- too much redundant and noisy data
  - E.g., reading just a few tweets (or news articles) is enough to know the majority of the topics in the stream.
- **Core Items:** informative or representative items in a data stream.
- **Core Items Tracking (CIT):** a streaming algorithm that can continuously track core items in a data stream in real-time.



# The Right to be Forgotten


- We want the CIT algorithm to be able to gradually forget historical data in the stream.




- CIT over insertion-only streams [KDD'14]:
- CIT over sliding-window streams [WWW'17]:

# The Right to be Forgotten

- We want the CIT algorithm to be able to gradually forget historical data in the stream.



- CIT over insertion-only streams [KDD'14]:



- CIT over sliding-window streams [WWW'17]:



# Outline

- Motivation ✓
- Problem formulation
- Algorithms
- Experiments
- Conclusion




# Measuring Informativeness of a Set of Items

- **Utility Function:**  $f: 2^V \mapsto \mathbb{R}_{\geq 0}$  where  $f(S)$  could measure the informativeness [JMLR'12, SIGIR'15], representativeness [KDD'14, CIKM'16, ICDE'18], diversity [WSDM'09], or coverage [PODS'14, KDD'15] of set  $S \subseteq V$ .
- $f(S)$  commonly satisfies the monotone submodular property, i.e.,
$$f(S \cup \{e\}) - f(S) \geq f(T \cup \{e\}) - f(T)$$
for all  $S \subseteq T \subseteq V$  and  $e \in V$ .
  - captures the **dimension return** property [Nemhauser et al. 1978]

# Probabilistic-Decaying Stream (PDS) Model

- At time  $t$ , we let an item  $e$  arrived at time  $t_e \leq t$  participate in the analysis with a probability  $p(e, t) = h_e(t - t_e)$ .
- $h_e: \mathbb{Z}_{\geq 0} \mapsto [0, 1]$  is an item-specific decaying function.
- $h_e(\text{age})$  decreases as  $\text{age}$  increases, e.g.,  $h_e(\text{age}) = p_e^{\text{age}}$ ,  $p_e \in (0, 1)$ .




# The Core Items Tracking (CIT) Problem

- **Given** a monotone submodular utility function  $f$ , a PDS with item-specific decaying function  $h_e$ , and a budget  $k > 0$
- **Want** to find a subset  $S_t^* \subseteq V$  at any query time  $t$ , s.t.

$$S_t^* = \arg \max_{S \subseteq V \wedge |S| \leq k} \mathbb{E}_{h_e} [f(S) | \mathcal{D}_t]$$

where  $\mathcal{D}_t \triangleq \{e: t_e \leq t\}$  denotes the items arrived before  $t$ .



prefer to choose more recent data items as core items

# Outline

- Motivation ✓
- Problem formulation ✓
- Algorithms
- Experiments
- Conclusion

# Expensive to Exactly Calculate $\mathbb{E}[f(S)|\mathcal{D}_t]$


- To calculate  $\mathbb{E}[f(S)|\mathcal{D}_t]$ , we have to consider the participation possibility of each item in  $S$ , e.g.,

$$\begin{aligned}\mathbb{E}[f(\{a, b\})|\mathcal{D}_t] = & \underbrace{p(a, t)p(b, t)f(\{a, b\})}_{\text{both } a \text{ and } b \text{ participate in the analysis}} \\ & + \underbrace{p(a, t)(1 - p(b, t))f(\{a\})}_{\text{only } a \text{ participates in the analysis}} + \underbrace{(1 - p(a, t))p(b, t)f(\{b\})}_{\text{only } b \text{ participates in the analysis}}.\end{aligned}$$

- Exactly calculating  $\mathbb{E}[f(S)|\mathcal{D}_t]$  requires  $O(2^{|S|})$  oracle calls.

# Monte-Carlo Approximation

- Generate  $n$  samples of the PDS, and estimate  $\mathbb{E}[f(S)|\mathcal{D}_t]$ .




- By Monte-Carlo approximation, we have

$$F(S) \triangleq \frac{1}{n} \sum_{i=1}^n f(S \cap \mathcal{D}_t^{(i)}) \xrightarrow{a.s.} \mathbb{E}[f(S)|\mathcal{D}_t], \quad n \rightarrow \infty.$$

- The number of oracle calls reduces from  $O(2^{|S|})$  to  $O(n)$ .
- $F(S)$  is still monotone and submodular.

# Overview




- **PND CIT** can efficiently solve the CIT problem over a special kind of probabilistic non-decaying case.
- **PDCIT** uses PND CIT as a building block to solve the CIT problem over general PDS.
- **PDCIT+** is designed to improve the efficiency of PDCIT.

algorithm	#oracle calls	memory	approx. ratio
PND CIT	$O(n\epsilon^{-1} \log k)$	$O(nk\epsilon^{-1} \log k)$	$1/2 - \epsilon$
PDCIT	$O(Ln\epsilon^{-1} \log k)$	$O(Lnk\epsilon^{-1} \log k)$	$1/2 - \epsilon$
PDCIT+	$O(n\epsilon^{-2} \log^2 k)$	$O(nk\epsilon^{-2} \log^2 k)$	$1/4 - \epsilon$

# PND CIT: Probabilistic Non-Decaying Case

- Probabilistic non-decaying case:  $p(e, t) \equiv p_e$
- The PDS can be converted to an insertion-only stream.




$$\begin{aligned} I(e_1) &= [0, 0, 0, 1]^T \\ I(e_2) &= [0, 1, 1, 1]^T \\ I(e_3) &= [1, 0, 1, 0]^T \end{aligned}$$

- $\{(e_1, p_{e_1}), (e_2, p_{e_2}), \dots\} \Rightarrow \{I(e_1), I(e_2), \dots\}$  where  $I(e) \in \{0, 1\}^n$
- submodular optimization over insertion-only streams has been extensively studied [SDM'08, DAM'12, SPAA'13, KDD'14].




# PDCIT: Probabilistic Decaying Case




$$t_e \quad l_0(e) = [1, 1, 1, 1]^T$$

↓


- At time  $t$ , denote those non-zero sample vectors by  $\mathcal{B}_t$ .
- Ideally, if we can feed  $\mathcal{B}_t$  to a PNDCIT instance, we will get a quality guaranteed solution at time  $t$ .
- How to process  $\mathcal{B}_t$  in a streaming fashion?




# PDCIT: Probabilistic Decaying Case



- At time  $t$ , denote those non-zero sample vectors by  $\mathcal{B}_t$ .
- Ideally, if we can feed  $\mathcal{B}_t$  to a PNDCIT instance, we will get a quality guaranteed solution at time  $t$ .
- How to process  $\mathcal{B}_t$  in a streaming fashion?




# PDCIT: Probabilistic Decaying Case




$$\begin{array}{l} t_e \\ t_e + 1 \\ \vdots \\ t_e + \infty \end{array} \left\{ \begin{array}{l} l_0(e) = [1, 1, 1, 1]^T \\ l_1(e) = [1, 1, 0, 1]^T \\ \vdots \\ l_\infty(e) = \mathbf{0} \end{array} \right.$$

- At time  $t$ , denote those non-zero sample vectors by  $\mathcal{B}_t$ .
- Ideally, if we can feed  $\mathcal{B}_t$  to a PNDCIT instance, we will get a quality guaranteed solution at time  $t$ .
- How to process  $\mathcal{B}_t$  in a streaming fashion?

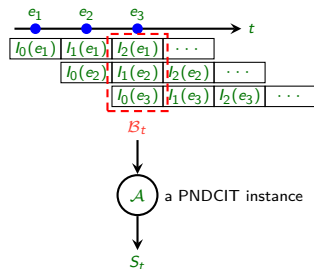


# PDCIT: Probabilistic Decaying Case




$$\begin{array}{c|c} t_e & l_0(e) = [1, 1, 1, 1]^T \\ t_e + 1 & l_1(e) = [1, 1, 0, 1]^T \\ \vdots & \vdots \\ t_e + \infty & l_\infty(e) = \mathbf{0} \end{array}$$

- At time  $t$ , denote those non-zero sample vectors by  $\mathcal{B}_t$ .
- Ideally, if we can feed  $\mathcal{B}_t$  to a PNDCIT instance, we will get a quality guaranteed solution at time  $t$ .
- How to process  $\mathcal{B}_t$  in a streaming fashion?




# PDCIT: Probabilistic Decaying Case

- Assume  $l_l(e) = \mathbf{0}$  for  $l \geq L, \forall e$ . Thus an item has at most  $L$  non-zero sample vectors  $\{l_l(e)\}_{0 \leq l < L}$ .
- PDCIT runs  $L$  PNDCIT instances, and processes each arrived item's sample vectors in parallel.
- At next time step, because each  $l_l(e)$  evolves to  $l_{l+1}(e)$ , which has been processed by PNDCIT instances on the right side, thus we shift these  $L$  PNDCIT instances one unit to the left.




# PDCIT: Probabilistic Decaying Case

- Assume  $l_l(e) = \mathbf{0}$  for  $l \geq L, \forall e$ . Thus an item has at most  $L$  non-zero sample vectors  $\{l_l(e)\}_{0 \leq l < L}$ .
- PDCIT runs  $L$  PNDCIT instances, and processes each arrived item's sample vectors in parallel.
- At next time step, because each  $l_l(e)$  evolves to  $l_{l+1}(e)$ , which has been processed by PNDCIT instances on the right side, thus we shift these  $L$  PNDCIT instances one unit to the left.





# PDCIT: Probabilistic Decaying Case

- Assume  $l_l(e) = \mathbf{0}$  for  $l \geq L, \forall e$ . Thus an item has at most  $L$  non-zero sample vectors  $\{l_l(e)\}_{0 \leq l < L}$ .
- PDCIT runs  $L$  PNDCIT instances, and processes each arrived item's sample vectors in parallel.
- At next time step, because each  $l_l(e)$  evolves to  $l_{l+1}(e)$ , which has been processed by PNDCIT instances on the right side, thus we shift these  $L$  PNDCIT instances one unit to the left.



# PDCIT: Probabilistic Decaying Case


- Assume  $l_l(e) = \mathbf{0}$  for  $l \geq L, \forall e$ . Thus an item has at most  $L$  non-zero sample vectors  $\{l_l(e)\}_{0 \leq l < L}$ .
- PDCIT runs  $L$  PNDCIT instances, and processes each arrived item's sample vectors in parallel.
- At next time step, because each  $l_l(e)$  evolves to  $l_{l+1}(e)$ , which has been processed by PNDCIT instances on the right side, thus we shift these  $L$  PNDCIT instances one unit to the left.





# PDCIT+: Improve Efficiency


- PDCIT needs to maintain  $L$  PNDCIT instances. What if  $L$  is extremely large?



- **Idea:** selectively maintain just a few PNDCIT instances, that can well approximate the rest
- Similar to using a histogram to approximate a curve

# PDCIT+: Improve Efficiency


- PDCIT needs to maintain  $L$  PNDCIT instances. What if  $L$  is extremely large?



- **Idea**: selectively maintain just a few PNDCIT instances, that can well approximate the rest
- Similar to using a histogram to approximate a curve

# PDCIT+: Improve Efficiency


- PDCIT needs to maintain  $L$  PNDCIT instances. What if  $L$  is extremely large?



- Idea:** selectively maintain just a few PNDCIT instances, that can well approximate the rest
- Similar to using a histogram to approximate a curve

# PDCIT+: Improve Efficiency


- PDCIT needs to maintain  $L$  PNDCIT instances. What if  $L$  is extremely large?




- **Idea**: selectively maintain just a few PNDCIT instances, that can well approximate the rest
- Similar to using a histogram to approximate a curve

# PDCIT+: Improve Efficiency

- PDCIT needs to maintain  $L$  PNDCIT instances. What if  $L$  is extremely large?



- Idea**: selectively maintain just a few PNDCIT instances, that can well approximate the rest
- Similar to using a histogram to approximate a curve



checkout our paper for more details

# Outline

- Motivation ✓
- Problem formulation ✓
- Algorithms ✓
- Experiments
- Conclusion

# Data

- **DBLP author stream**: an item is a set of conferences that the author attended before
- **MemeTracker article stream**: an item is a set of memes that the article contains
- **math.StackExchange question stream**: an item is a set of tags that the question contains
- **StackOverflow question stream**: an item is a set of tags that the question contains

data stream	item	length	time period
DBLP	author	371,690	1936 - 2018
MemeTracker	article	714,072	1/2009 (one month)
math.StackExchange	question	955,284	7/2010 - 6/2018
StackOverflow	question	2,904,450	1/2015 - 3/2016



# Settings

- **Goal:** maintain  $k$  most representative items that jointly have the maximum coverage, i.e.,  $f(S) = |\cup_{e \in S} e|$ .
- **Decaying Function:**  $h_e(x) = p^x, p \in (0, 1)$ .
- **Baselines:**
  - GREEDY will serve as an upper bound
  - Unbiased Reservoir Sampling
  - Biased Reservoir Sampling



# PDCIT vs PDCIT+

- How close is their solution quality?
- How significant can PDCIT+ reduce the number of oracle calls?

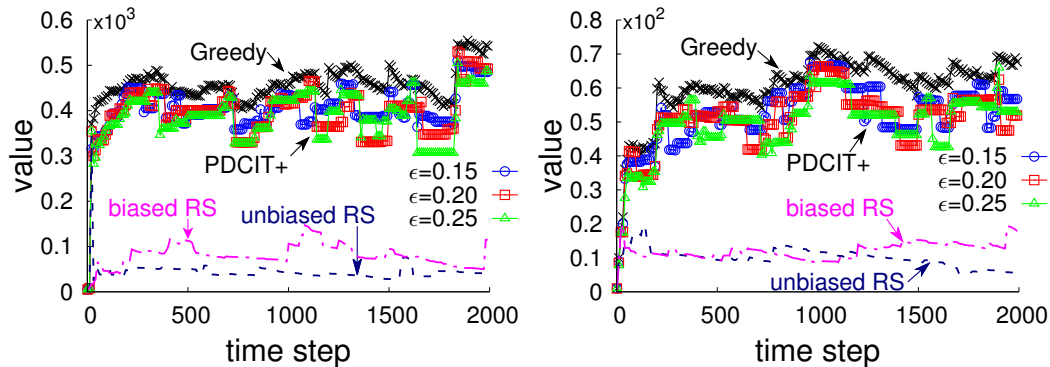


left: DBLP, right: MemeTracker,  $\epsilon = 0.2, k = 10, n = 20, L = 100$

**achieves similar solution quality, reduces more than a half of oracle calls**

# Solution Quality



- Comparing solution quality of different methods (higher is better)



left: DBLP, right: MemeTracker,  $p = 0.999$ ,  $k = 10$ ,  $n = 20$

# Scalability

- Comparing the number of oracle calls of different methods (lower is better)



left: DBLP, right: MemeTracker,  $p = 0.999$ ,  $k = 10$ ,  $n = 20$

# Outline

- Motivation ✓
- Problem formulation ✓
- Algorithms ✓
- Experiments ✓
- Conclusion

# Conclusion

- The optimization problem behind the CIT problem is to solve the streaming submodular optimization problem over probabilistic-decaying streams.
- We designed two streaming algorithms, namely PDCIT and PDCIT+, to address this CIT problem. They use PNDCIT as a building block.
- These techniques are verified on several public available data streams. The results demonstrate the effectiveness of these methods.

Thanks for listening!

[junzhou.zhao@xjtu.edu.cn](mailto:junzhou.zhao@xjtu.edu.cn)